

ELEC3608 Computer Architecture

Instruction Set Architecture

Philip H.W. Leong
School of Electrical and Information Engineering,
The University of Sydney



THE UNIVERSITY OF
SYDNEY

Introduction

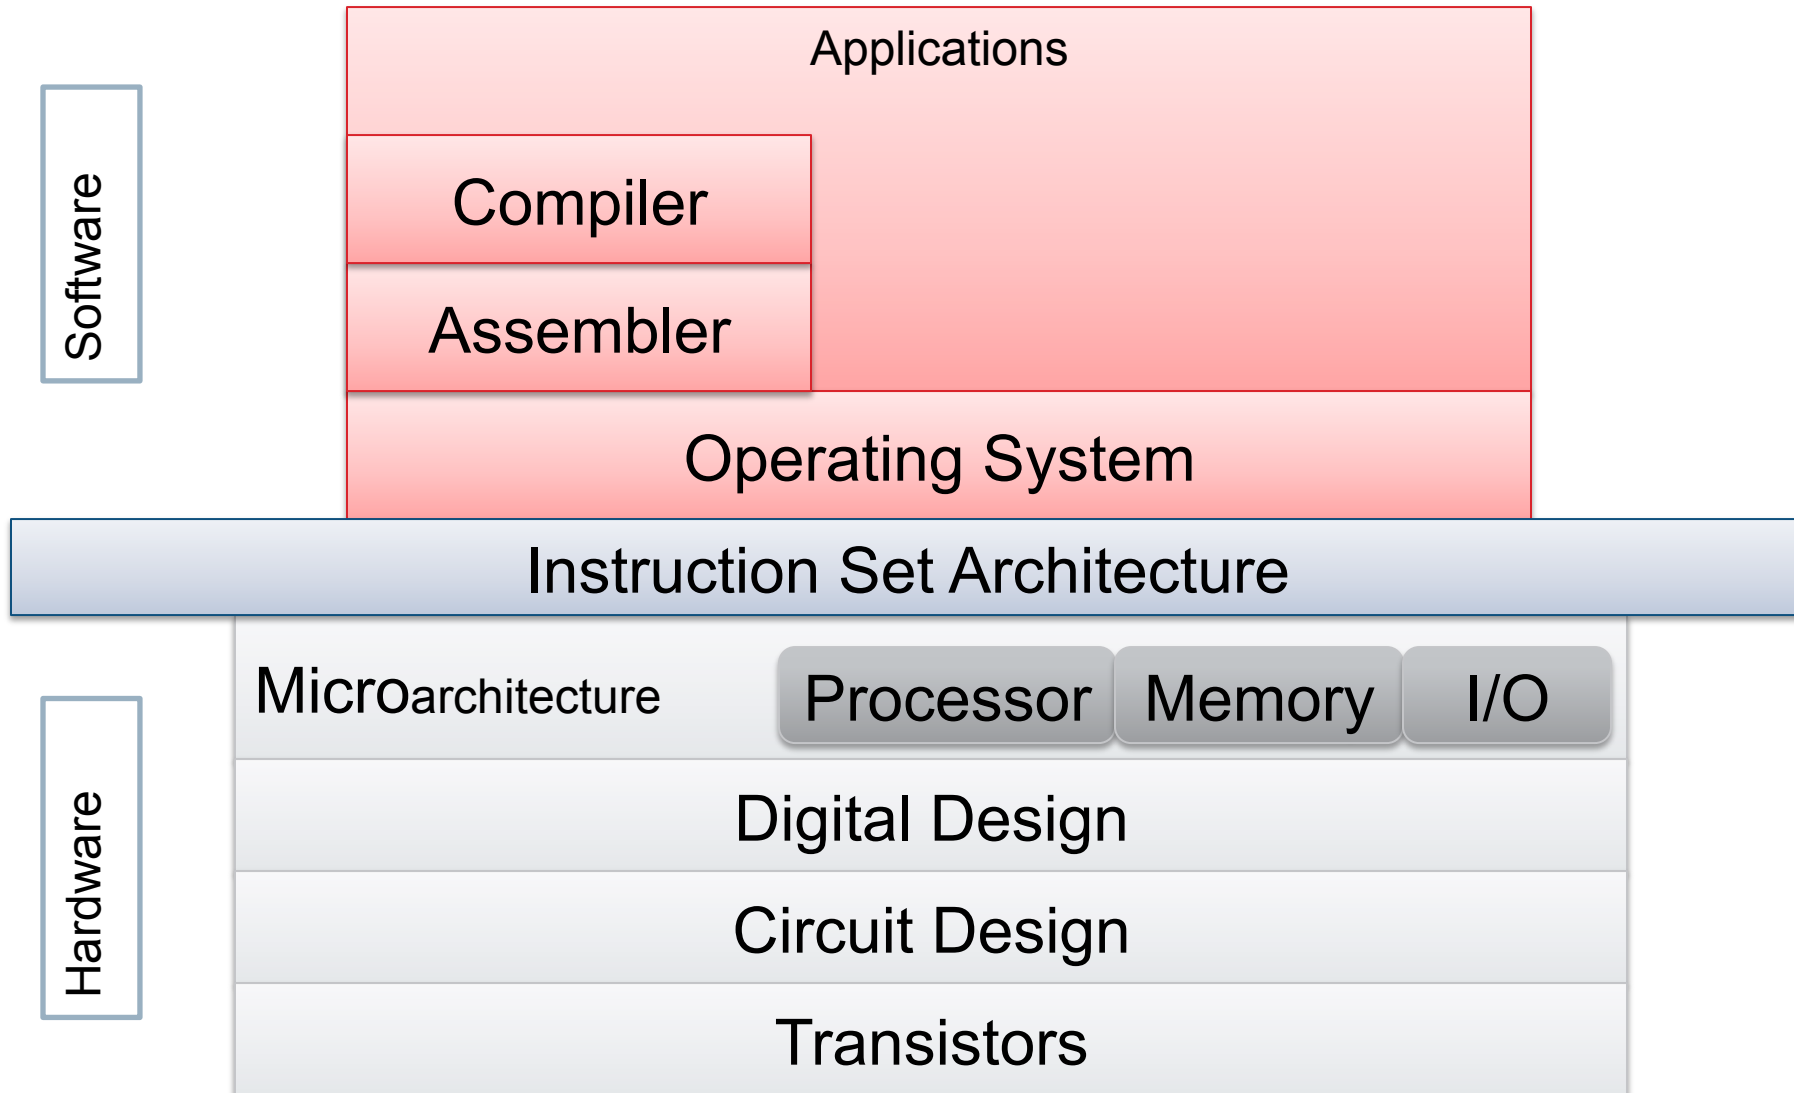


THE UNIVERSITY OF
SYDNEY

- › Instruction set architecture defines the user observable behavior of a processor
 - A contract between hardware and software
- › Usually includes:
 - Observable state of a processor
 - A set of machine instructions
 - Semantics of the instruction and processor execution

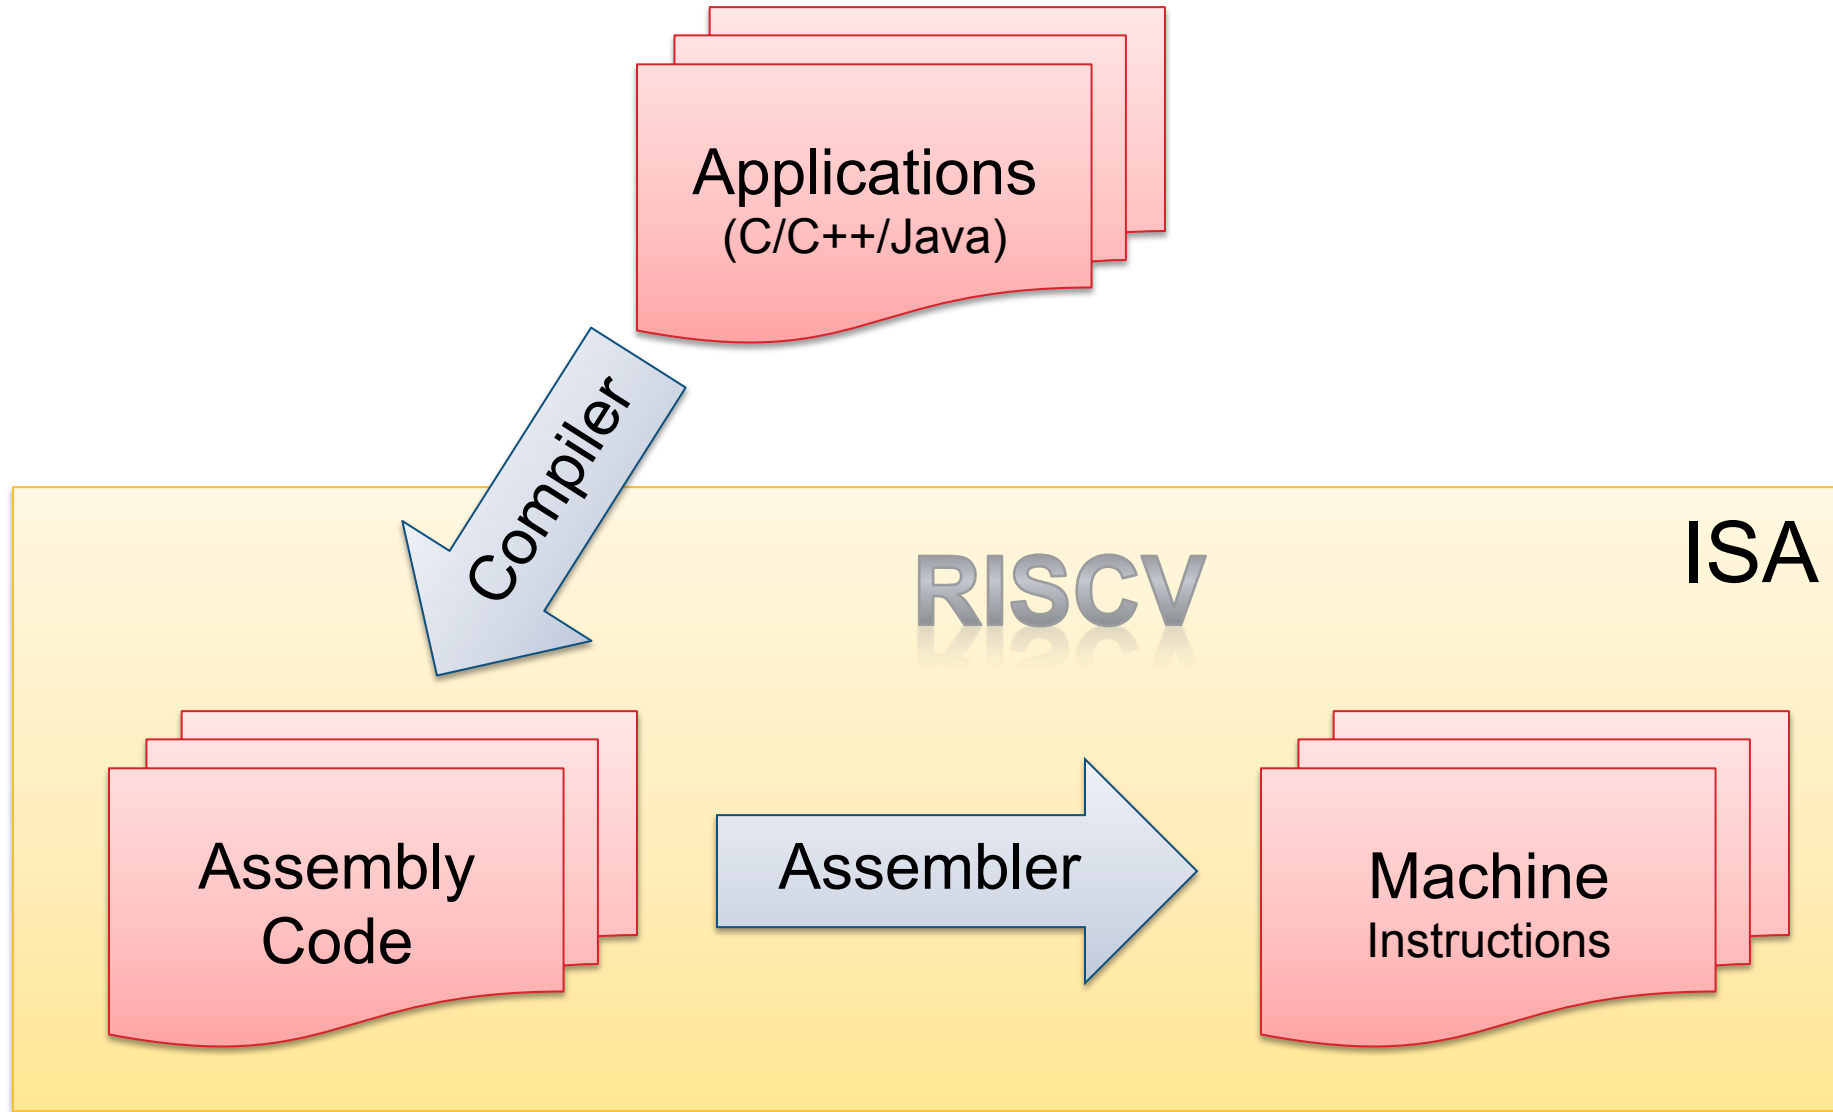


Computer Architecture: HW/SW Interface





Review: Compiling Applications



Note about learning an ISA

- › By learning how ONE ISA is designed, you learn how various architectural tradeoffs are performed.
- › NOT to memorize details of a single ISA
- › The subset of RISC-V covered in class is very similar to MIPS:
 - More examples in COD
 - Will use MIPS software for simulation in some cases

RISC-V ISA Version 2.0

- › New RISC design from UC Berkeley
- › Realistic & complete ISA, but open & small
- › Not over-architected for a certain implementation style
- › 32-bit and 64-bit address space variants
 - RV32 and RV64
- › Designed to be extensible
 - High-performance multiprocessing
 - Efficient instruction encoding for embedded system
 - etc
- › Easy to subset/extend for education/research
- › Complete compilation toolchain
- › FPGA and ASIC designs



<http://www.riscv.org>

- › A revised version to the original definition
- › Divides into an integer base ISA + extension
 - Floating point, etc
- › Revised instruction coding
 - Make hardware design easier
- › Added 128-bit support
- › Much more...
- › We will be using RISC-V ISA 2.0 this semester:
 - Things may break
 - RISC-V base integer ISA is similar to MIPS in textbook
 - Will try cross reference between RISCV and MIPS

- › The riscv processor has a **load-store architecture**
 - Most instructions operate on values directly to and from the **register file**
 - Dedicated **load/store instructions** to transfer data between register file and main memory
- › RV32 has a 32-bit **word length**
 - Fixed instruction length of 32 bits
 - A data word is 32 bits wide
 - RISC-V ISA designed to allow variable length instruction code and instruction extension
 - Not covered in this course

RV32 Processor State – Base Integer Design

- › Register file with **32** 32-bit general-purpose registers
- › Labeled x0 – x31
 - x0 always contains the value 0
- › Application Binary Interface (ABI) names used in assembler
 - t0 – t6 for temp variables
 - s0 – s11 for saved variables, etc
- › Additional **Program Counter** (PC) register
 - Contains the memory address of the current executing instruction

Register	ABI Name
x0	zero
x1	ra
x2	sp
x3	gp
x4	tp
x5-7	t0-2
x8	s0/fp
x9	s1
x10-17	a0-7
x18-27	s2-11
x28-31	t3-6

XLEN=32 in RV32

<https://blog.riscv.org/2015/01/announcing-the-risc-v-gcc-4-9-port-and-new-abi/>

Integer Register-Register Operations

- › Basic arithmetic and bit operations:
 - Arithmetic: ADD, SUB
 - Bit Operations: AND, OR, XOR, SLL, SRL, SRA
 - Comparison: SLT, SLTU
- › Take two **source** register inputs (denoted rs1, rs2) and produce one **destination** register output (denoted rd).



Assembler	Semantics
ADD rd, rs1, rs2	$rd \leftarrow rs1 + rs2$
SUB rd, rs1, rs2	$rd \leftarrow rs1 - rs2$

- › 3 operands:
 - 2 sources and 1 destination
- › Add/sub the values of the 2 source registers (rs1, rs2) and store the result in destination register (rd)
- › Discard overflow
- › Lower 32 bits are written

› C code:

```
y = (e + f) - (g + h);
```

› Assume e, f, g, h, y stored in a0, a1, a2, a3, a4 respectively.

› RISC-V code:

```
ADD t0, a0, a1  
ADD t1, a2, a3  
SUB a4, t0, t1
```

Assembler	Semantics
AND rd, rs1, rs2	$rd \leftarrow rs1 \text{ AND } rs2$
OR rd, rs1, rs2	$rd \leftarrow rs1 \text{ OR } rs2$
XOR rd, rs1, rs2	$rd \leftarrow rs1 \text{ XOR } rs2$

- › Similar semantics to ADD/SUB instructions
- › Bit operations useful for masking values
- › E.g.:

```
# t0 ← 0xABCD0123
# t1 ← 0x000F0000
AND t2, t0, t1
# t2 contains 0x000D0000
```




```
# t0 ← 0x00C00005  
# t1 ← 0xFFFFFFFF  
XOR t2, t0, t1
```

What is the value stored in t2?

Instruction Encoding

RISC-V Instruction Encoding

- › Each supported instruction in the ISA must be uniquely encoded to describe its function
- › Any unique encoding works, but good encoding may optimize:
 - Regularity; Simple hardware decode; Instruction length; Extensibility
- › In RV32 6 general types of instructions: R, I, S, SB, U, UJ
 - Favors hardware regularity over human readability
 - Each instruction is 32-bit wide

31	30	25	24	21	20	19	15	14	12	11	8	7	6	0				
funct7				rs2			rs1		funct3		rd			opcode		R-type		
imm[11:0]						rs1		funct3		rd			opcode		I-type			
imm[11:5]				rs2			rs1		funct3		imm[4:0]			opcode		S-type		
imm[12]		imm[10:5]			rs2			rs1		funct3		imm[4:1]		imm[11]		opcode		SB-type
imm[31:12]										rd			opcode			U-type		
imm[20]		imm[10:1]			imm[11]		imm[19:12]				rd			opcode		UJ-type		

R-type Instruction Encoding

- › All the simple integer register-register operations mentioned so far are encoded as R-type instructions

R-Type format

31	25 24	20 19	15 14	12 11	7 6	0
funct7	rs2	rs1	funct3	rd	opcode	
7	5	5	3	5	7	
0000000	src2	src1	ADD/SLT/SLTU	dest	OP	
0000000	src2	src1	AND/OR/XOR	dest	OP	
0000000	src2	src1	SLL/SRL	dest	OP	
0100000	src2	src1	SUB/SRA	dest	OP	

- › Instructions within the same type shares similar encoding
- › Differentiate only by the funct3 and funct7 fields



- › If the size of register file is increased to include 64 registers, how many bits are needed to encode rs1?
- › How would the above change affect the number of instructions representable?

- › Very often programs need to operate on (small) numerical constants in the code
 - e.g.: “a = b + 3”
- › RV32 includes register-immediate instructions for these cases
- › Small constant embedded within the instruction.
 - Arithmetic: ADDI
 - Bit Operations: ANDI, ORI, XORI
 - Comparison: SLTI, SLTIU
- › Take one **source** register inputs (rs1) and produce one **destination** register output (rd).

I-type Instruction Encoding

I-Type format

- › Most register-immediate operations are encoded using the I-type format

31	20 19	15 14	12 11	7 6	0
imm[11:0]		rs1	funct3	rd	opcode
12		5	3	5	7
I-immediate[11:0]		src	ADDI/SLTI[U]	dest	OP-IMM
I-immediate[11:0]		src	ANDI/ORI/XORI	dest	OP-IMM

- › NOTE: since the constant is encoded *within* the instruction, size of immediate must be < 32 bits
- › In RV32, I-type instruction, immediate constants are encoded using 12 bits
- › Note that bit locations of rs1, funct3, rd and opcode are the same as in R-type
 - Simplicity → high performance

› C code:

```
y = y + e + 1
```

› Assume e, y stored in a0, a1 respectively.

› RISC-V code:

```
ADD  t0, a1, a0  
ADDI a1, t0, 1
```

Why is there no SUBI instruction... ?

Number Systems

Unsigned Numbers


- › Represent non-negative binary numbers (0, 1, 2, 3, ...) using their natural binary representations

- › An n-bit bitstring can represent numbers in

$$[0, 2^n - 1]$$

- › Represents **equally spaced** integers on the number line

Value	Binary	Bitstring (8-bit)
0	0	00000000
1	1	00000001
2	10	00000010
3	11	00000011
4	100	00000100
5	101	00000101
6	110	00000110
7	111	00000111
8	1000	00001000
9	1001	00001001
10	1010	00001010
11	1011	00001011



Two's Complement (1)

- › Negative numbers are represented by the 2's complement of the absolute value of that number
 - 2's complement of an n -bit number v is the value $2^n - v$
- › For example, to represent value -3 using a 4-bit bitstring:
 - $3 \rightarrow "0011"$
 - $2^4 - 3 = 16 - 3 = 13$
 - $\therefore -3 \rightarrow "1101"$

- › 2's complement of a number can be obtained by “adding 1 to the 1's complement of the number”

Original	23	00010111
1's complement		11101000
2's complement	-23	11101001

- › The value of a bitstring $\{b_{n-1}b_{n-2}\cdots b_0\}$ in 2's complement can be calculated as:

$$-2^{n-1} + \sum_{i=0}^{n-2} 2^i b_i$$

- › Range: $[-2^{n-1}, 2^{n-1} - 1]$
 - Note that is asymmetric

› What is the binary representation of 13?

1101_2

› What is the binary representation of -13?

1

0010

2

0011

3

Cannot be represented

4

Cannot be determined based on given information

› What are the binary representations of 13 and -13 as 8-bit integers?

1

13: 00001101 -13: 00000011

2

13: 00001101 -13: 11110011

› What are the binary representations of 13 and -13 as 12-bit integers?

1

13: 000000001101 -13: 000011110011

2

13: 000000001101 -13: 111111110011


- › The sign of an integer can only be known with the width of integer specified
- › In 2's complement representations, the most significant bit (MSB) of an integer denotes its sign:
 - MSB=1 → negative
 - MSB=0 → non-negative
- › To cast an integer value to a representation with more bits, it must be **sign-extended** to preserve its sign:
 - If it is negative number, fill the additional bits with 1
 - If it is non-negative, fill the additional bits with 0

ADDI with negative numbers

31	20 19	15 14	12 11	7 6	0
imm[11:0]	rs1	funct3	rd	opcode	
12	5	3	5	7	
I-immediate[11:0]	src	ADDI/SLTI[U]	dest	OP-IMM	
I-immediate[11:0]	src	ANDI/ORI/XORI	dest	OP-IMM	

- › Subtraction with immediate values can be achieved by using negative immediate with ADDI
- › Recall that immediate constants are 12 bits long in all I-type instructions
 - Range: $[-2^{11}, 2^{11}-1]$
- › In general, instructions must clearly specify in its semantics if it is treating the value as signed or unsigned numbers
 - E.g. SLTI vs SLTIU

Load Upper Immediate



- › Load upper bits of immediate constant
- › Form constants of any values in combination with ADDI, etc
- › No need for sign extension
- › NOTE: opcode and rd in the same location as most other instructions

Memory Operations

- › All program data + instruction originally stored in main memory
 - Array, structure, user data, string, etc
- › Dedicated instruction to load/store data between register file and main memory
- › NOTE: registers are much faster than main memory
 - Limited registers availability (32 registers vs. 4GB of memory)
 - Compiler must use registers wisely for best performance

- › Memory is byte-addressable
 - Each address corresponds to 1 byte of data
- › Native data types:
 - Word: 32 bit (`int`)
 - Half-word: 16 bit (`short`)
 - Byte: 8 bit (`char`)
- › RV32 has a **Little-Endian** byte addressing scheme
 - The least significant byte is at the lowest address



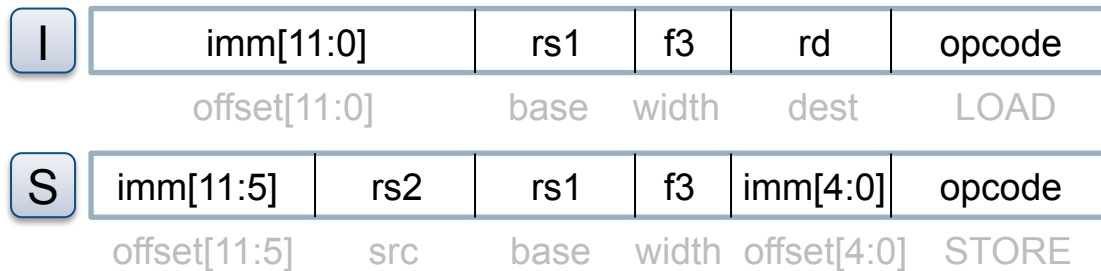
```
int A[20]  
A[0]=0xABCD0123  
A[1]=0xFACECAFE
```

	Big Endian	Little Endian
0x000A0007	FE	FA
0x000A0006	CA	CE
0x000A0005	CE	CA
0x000A0004	FA	FE
0x000A0003	23	AB
0x000A0002	01	CD
0x000A0001	CD	01
0x000A0000	AB	23

- › Words are **naturally aligned** at 4-byte boundary
- › Half-words naturally aligned at 2-byte boundary
- › RV32I allows non-aligned access
 - **Not common** in most processors
 - See ISA spec

		aligned	aligned	Not aligned
0x000A0007			AB	
0x000A0006			CD	
0x000A0005			01	AB
0x000A0004			23	CD
0x000A0003	AB			01
0x000A0002	CD			23
0x000A0001	01			
0x000A0000	23			

Memory Operations



Load:

$(\text{dest}) \leftarrow M[(\text{base}) + \text{offset}]$

Store:

$M[(\text{base}) + \text{offset}] \leftarrow (\text{src})$

- › Immediate is sign extended
- › LW/SW – Word
- › LH/SH – Half Word
- › LB/SB – Byte

```
int y = e + f[3];
```

- › Assume e, f, y stored in a0, a1, a2 respectively.
- › RISC-V code:

```
lw    t0, 12(a1)    # 12 = 3 x 4  
ADD   a2, a0, t0
```

- › LW, LH, LB instructions load word, half-word and byte from memory

Memory Operations – Store

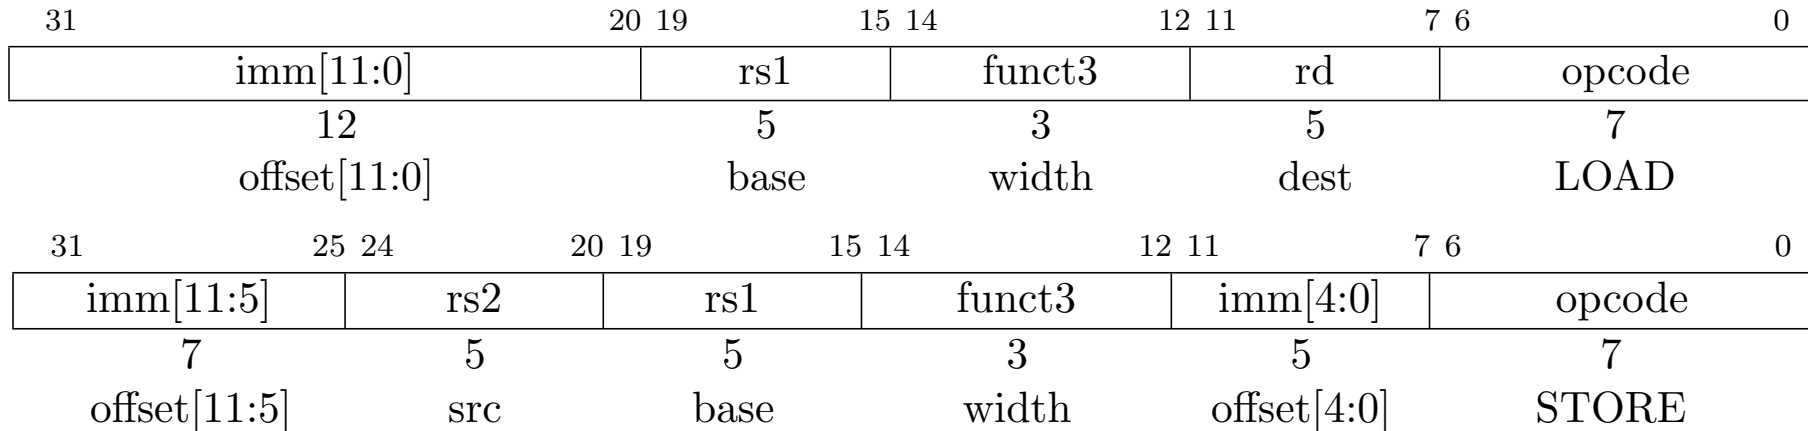
```
int f[7] = f[6] + f[5];
```

- › Assume f is stored in a0.
- › RISC-V code:

```
lw  t0, 20(a0)    # 20 = 5 x 4
lw  t1, 24(a0)
ADD t0, t0, t1
sw  t0, 28(a0)
```

- › SW, SH, SB instructions store word, half-word and byte to memory

Load/Store Instruction Encoding



- › LOAD instructions are encoded as I-type
- › STORE instructions are encoded as S-type
- › funct3 encodes the length (W, H, B)
- › 12-bit offset is sign extended and added to base address in rs1
- › Note the bit location of imm[11:5] and imm[4:0] in S-type



Immediate Formats

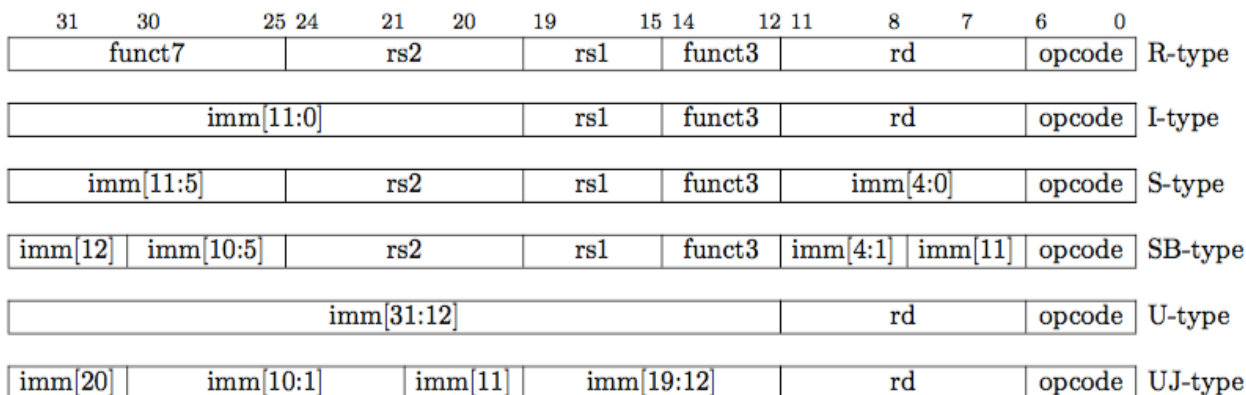


Figure 2.3: RISC-V base instruction formats showing immediate variants.

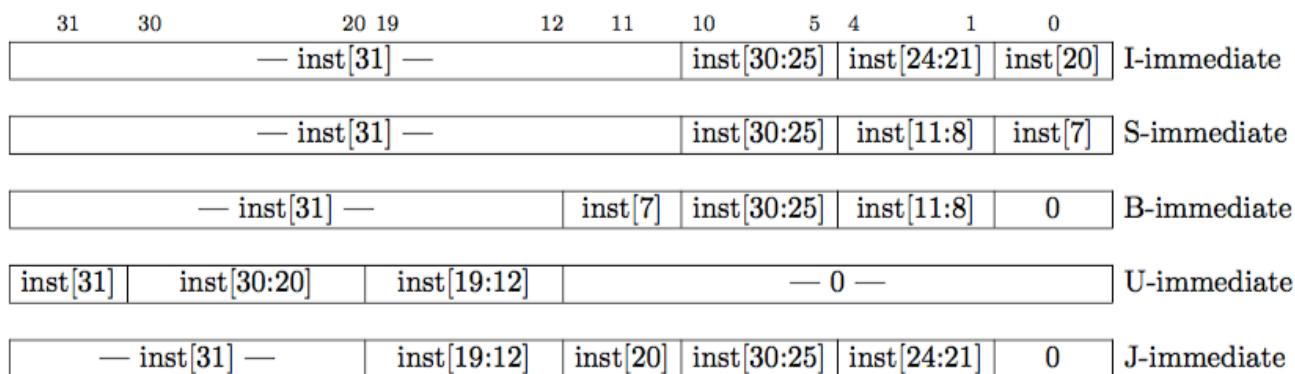


Figure 2.4: Types of immediate produced by RISC-V instructions. The fields are labeled with the instruction bits used to construct their value. Sign extension always uses inst[31].

Branch Operations

- › Normally programs execute sequentially
 - The next instruction to fetch is normally at location $PC + 4$
- › Decision making constructs requires fetching instruction from run-time determined location
 - if-the-else, do...while, goto, switch, etc
- › RISC-V has two types of control transfer instructions:
 - **Unconditional jump**
 - **Conditional branch**



Conditional Branches

SB-Type format


31				30				25 24				20 19				15 14				12 11				8				7				6				0			
imm[12]				imm[10:5]				rs2				rs1				funct3				imm[4:1]				imm[11]				opcode											
1				6				5				5				3				4				1				7											
offset[12,10:5]								src2				src1				BEQ/BNE								offset[11,4:1]								BRANCH							
offset[12,10:5]								src2				src1				BLT[U]								offset[11,4:1]								BRANCH							
offset[12,10:5]								src2				src1				BGE[U]								offset[11,4:1]								BRANCH							

› 12-bit branch immediate encoded as *signed offset in multiples of 2*

- E.g. offset[12:1] = 0x00C → jumps to PC + 12*2 = PC + 24

Assembler	Semantics
BEQ rs1, rs2, dest	if rs1=rs2 then branch dest
BNE rs1, rs2, dest	If rs1≠rs2 then branch dest
BLT rs1, rs2, dest	If rs1<rs2 then branch dest
BGE rs1, rs2, dest	If rs1≥rs2 then branch dest
BLTU, BGEU	Unsigned comparison of BLT, BGE

Unconditional Jump-and-Link



UJ-Type format

› Semantics:

- Unconditional jump to $PC + \text{offset} * 2$
- Store return address ($PC + 4$) in rd

Jump

- › J-immediate encodes a signed offset in multiples of 2 bytes
 - 20 bits → $\pm 1\text{MiB}$ range
- › Jump relative to PC
- › Plain jump (J assembler) use rd=x0

- › To facilitate function calls
 - Call to function and return to the next instruction
- › Convention is to use x1 (ra) as return register

```
v = 9  
r = sqrt(v)  
r = r + 27
```


```
0xA000 addi a0, zero, 9  
0xA004 jal  ra, sqrt  
0xA008 addi a0, a0, 27  
.  
.  
0xBC04 <. . .>  
0xBC08 <. . .>  
0xBC0C <. . .>  
0xBC10 return
```

+ 0x1C00

- return to 0xA008
- 0xA008 stored in register ra from the jal instruction

Unconditional JALR

I-Type format



- › Similar to JAL, except destination is relative to the value of rs1
- › Semantics:
 - Unconditional jump to offset + (rs1)
 - Store return address (PC+4) in rd
- › Note
 - the offset is not offset by multiple of 2
 - Entirely same as an I-type format instruction



JAL + JALR for function calls

```
0xA000 addi a0, zero, 9
0xA004 jal  ra, 0x1C00
0xA008 addi a0, a0, 27
.
.
0xBC04 <. . .>
0xBC08 <. . .>
0xBC0C <. . .>
0xBC10 jalr zero, ra
```



If-Then-Else Conversion

```
if (a >= b) {  
    c = a - b  
} else {  
    c = a + b  
}
```

```
        blt a1, a0, truepart  
        add  a2, a1, a0  
        j    exit  
truepart: sub a2, a1, a0  
exit    :  <...>
```

Procedure Calls

- › Steps required
 1. Place parameters in registers
 2. Transfer control to procedure
 3. Acquire storage for procedure
 4. Perform procedure's operations
 5. Place result in register for caller
 6. Return to place of call
-



Register Usage

Register	ABI Name	Description	Saver
x0	zero	Hard-wired zero	—
x1	ra	Return address	Caller
x2	s0/fp	Saved register/frame pointer	Callee
x3–13	s1–11	Saved registers	Callee
x14	sp	Stack pointer	Callee
x15	tp	Thread pointer	Callee
x16–17	v0–1	Return values	Caller
x18–25	a0–7	Function arguments	Caller
x26–30	t0–4	Temporaries	Caller
x31	gp	Global pointer	—
f0–15	fs0–15	FP saved registers	Callee
f16–17	fv0–1	FP return values	Caller
f18–25	fa0–7	FP arguments	Caller
f26–31	ft0–5	FP temporaries	Caller

› C code:

```
int leaf_example (int g, h, i, j)
{ int f;
  f = (g + h) - (i + j);
  return f;
}
```

- Arguments g, ..., j in a0, ..., a3
 - f in s0 (hence, need to save s0 on stack)
 - Result in a0
-

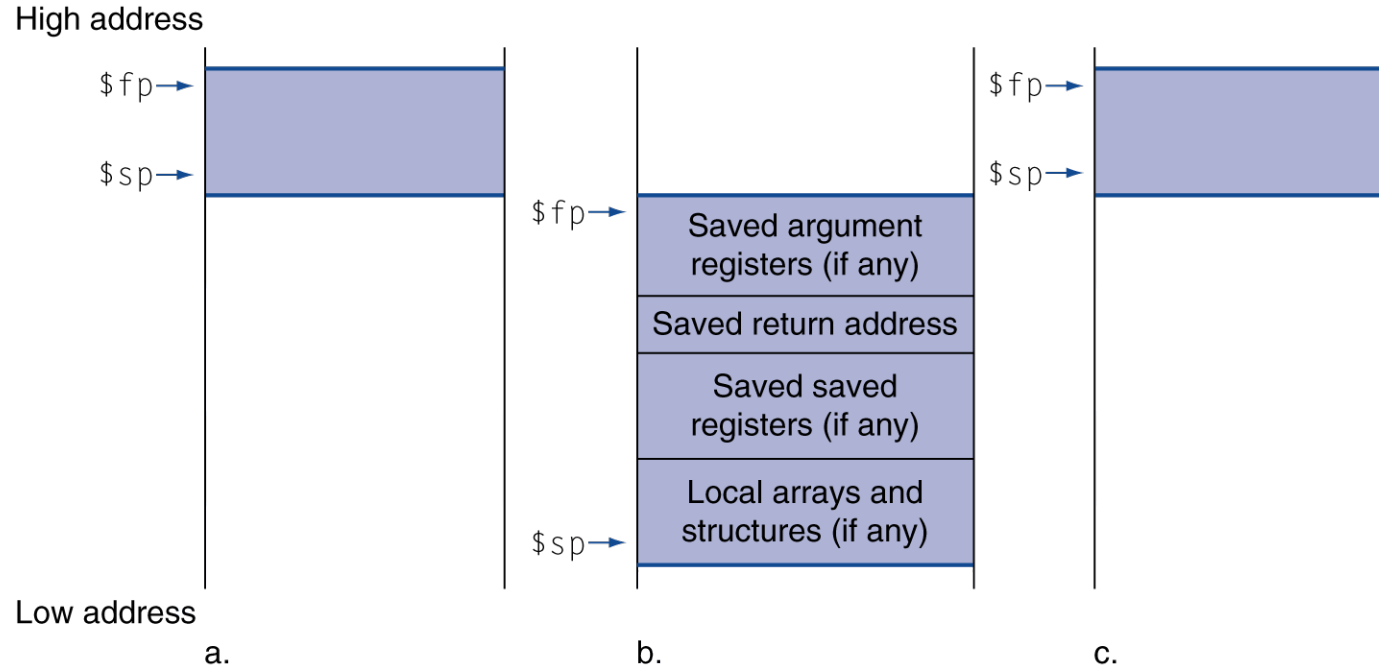
› RISC-V code:

`leaf_example:`

<code>addi sp, sp, -4</code>	Save s0 on stack
<code>sw s0, 0(sp)</code>	
<code>add t0, a0, a1</code>	Procedure body
<code>add t1, a2, a3</code>	
<code>sub s0, t0, t1</code>	
<code>add a0, s0, zero</code>	Result
<code>lw s0, 0(sp)</code>	Restore s0
<code>addi sp, sp, 4</code>	
<code>jalr zero, ra</code>	Return




Local Data on the Stack



- › Local data allocated by callee
 - e.g., C automatic variables
- › Procedure frame (activation record)
 - Used by some compilers to manage stack storage

- › Text: program code
- › Static data: global variables
 - e.g., static variables in C, constant arrays and strings
 - \$gp initialized to address allowing \pm offsets into this segment
- › Dynamic data: heap
 - E.g., malloc in C, new in Java
- › Stack: automatic storage



- › Procedures that call other procedures
 - › For nested call, caller needs to save on the stack:
 - Its return address
 - Any arguments and temporaries needed after the call
 - › Restore from the stack after the call
-

› C code:

```
int fact (int n)
{
    if (n < 1) return 1;
    else return n * fact(n - 1);
}
```

- Argument n in a0
 - Result in a0
-

› RISC-V code:

fact:

```
    addi sp, sp, -8      # adjust stack for 2 items
    sw   ra, 4(sp)       # save return address
    sw   a0, 0(sp)       # save argument - local variable
    li   t0, 2           # put a 2 into t0
    bgeu a0, t0, L1      # if n >= 2 take jump
    li   a0, 1           # otherwise, result is 1
    addi sp, sp, 8       # pop 2 items from stack
    jr   ra              # and return
L1:  addi a0, a0, -1      # else decrement n
     jal fact            # recursive call to get fact(n-1)
     lw   t0, 0(sp)      # restore original n
     mul  a0, t0, a0      # multiply to get result n * fact(n-1)
     lw   ra, 4(sp)      # restore return address
     addi sp, sp, 8      # pop 2 items from stack
     jr   ra             # and return
```

- › RISC-V ISA is a new academic RISC ISA used as examples
 - RV32 forms the base set of architectural feature
 - 32-bit architecture
 - 32 General Purpose Register File
- › Each instruction in an ISA must be uniquely encoded.
 - Encoding and semantics of ISA must tradeoff between hardware efficiency and usability.

- › These slides contain material developed and copyright by:
 - Arvind (MIT)
 - Krste Asanovic (MIT/UCB)
 - Joel Emer (Intel/MIT)
 - James Hoe (CMU)
 - John Kubiatowicz (UCB)
 - David Patterson (UCB)
 - Hayden So (HKU)
- › MIT material derived from course 6.823
- › UCB material derived from course CS152, CS252
- › HKU material derived from course ELEC3441